

nova  
net

# Caching og K6

Hans Arne Vartdal

Lars Alexander Jakobsen

Utvikler på **platform**teamet med fokus på **API**.

Brukes av Smartbokser, Smart TV, AppleTV,  
AndroidTV, iOS, Android, og web.

**RiksTV**

Noen tall fra siste 7 dager:

**94,5 millioner** requests

**~800** requests per second (peak)

Snitt responstid **38,9ms**

**1,8GB** komprimert JSON


**RiksTV**



# Ingen cache


Trenger du egentlig cache?

Systemer uten cache er **lette å forstå**, og  
**lette å teste**.




# ConcurrentDictionary

- Lett å implementere, og lett å teste
- Lite “out of the box”
- Må implementere strategier for oppdatering/invalidering
- Relativt små datasett med lav endringstakt




# MemoryCache

- Bruk **Microsoft.Extensions.Caching.Memory**
- Godt egnet for «cache aside» pga støtte for **Time To Live** (TTL)
- Lar deg konfigurere **minneforbruket**
- Lar deg håndtere når data kastes ut (cache eviction)




# LazyCache

- Open source bibliotek som **utvider MemoryCache**
- Garanterer at oppdatering av cache bare kjøres en gang
- **Unngå flere samtidige kall** mot en treg tjeneste






nova  
net



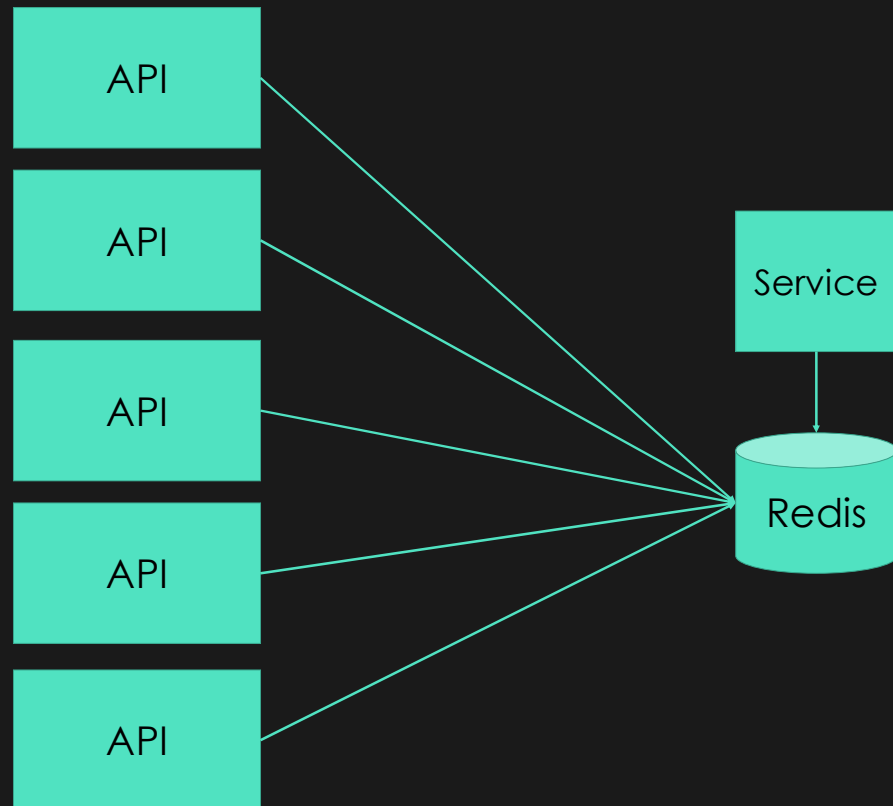
# Distribuert cache

- Distribuert cache; Ikke i samme prosess som applikasjonen
- For eksempel **Redis**
- Det finnes flere strategier for å oppdatere en slik cache!

# Lazy loading




# Write through




nova  
net

# Lazy loading + Write through




nova  
net

**In memory**  
**Lazy cache**  
**+**  
**Redis**  
**Lazy loading**  
**+**  
**Redis**  
**Write through**



nova  
net

# Event driven cache



# Ting å tenke på

- Må **første request** være like rask?
- **Hvor mye** data har jeg?
- **Hvor ofte** må jeg oppdatere dataene?
- Kaldstart
- **Versjonering** av distribuert cache?
- Distribuert cache kan bli «single point of failure»
- Tenk på **latency** ved distribuert cache.
- **Immutability** for InMemory cache



nova  
net

# K6

- Verktøy for ytelsestesting
  - Tester i Javascript
  - Command Line verktøy

Linux:

```
sudo apt-get install k6
```

Mac:

```
brew install k6
```

Windows: <https://dl.bintray.com/loadimpact/windows/k6-v0.26.1-amd64.msi>

Docker:

```
docker pull loadimpact/k6
```



nova  
net

# K6

- Open source, gratis og laget for utviklere
- Lett å komme i gang
  - Kan starte med ytelsestesting fra utviklermaskin
- Kan integreres i CI/CD

nova  
net

DEMO

# Enkel test

```
import http from 'k6/http';  
import { sleep } from 'k6';  
  
export default function() {  
  
    let result = http.get('https://novanetk6api.azurewebsites.net/nocache/99');  
  
    sleep(1); //sekunder  
  
}
```

# Enkel test: Resultat



```
execution: local-  
  output: -  
  script: ./scripts_01.js
```

```
duration: -, iterations: 1  
  vus: 1, max: 1
```

```
done [=====] 1 / 1i
```

```
data_received.....: 6.7 kB 2.8 kB/s  
data_sent.....: 565 B 240 B/s  
http_req_blocked.....: avg=292.92ms min=292.92ms med=292.92ms max=292.92ms p(90)=292.92ms p(95)=292.92ms  
http_req_connecting.....: avg=38.99ms min=38.99ms med=38.99ms max=38.99ms p(90)=38.99ms p(95)=38.99ms  
http_req_duration.....: avg=1.05s min=1.05s med=1.05s max=1.05s p(90)=1.05s p(95)=1.05s  
http_req_receiving.....: avg=0s min=0s med=0s max=0s p(90)=0s p(95)=0s  
http_req_sending.....: avg=0s min=0s med=0s max=0s p(90)=0s p(95)=0s  
http_req_tls_handshaking...: avg=133.96ms min=133.96ms med=133.96ms max=133.96ms p(90)=133.96ms p(95)=133.96ms  
http_req_waiting.....: avg=1.05s min=1.05s med=1.05s max=1.05s p(90)=1.05s p(95)=1.05s  
http_reqs.....: 1 0.425313/s  
iteration_duration.....: avg=2.35s min=2.35s med=2.35s max=2.35s p(90)=2.35s p(95)=2.35s  
iterations.....: 1 0.425313/s  
vus.....: 1 min=1 max=1  
vus_max.....: 1 min=1 max=1
```

```
PS D:\dev\novanet\k6-caching\Novanet.K6.Api\tests> |
```

# Options

```
import http from 'k6/http';
import { sleep } from 'k6';

export let options = {
  vus: 5,
  duration: "30s"
};

export default function() {

  let result = http.get('https://novanetk6api.azurewebsites.net/nocache/99');

  sleep(1); //sekunder
}
```

# Checks

```
import http from 'k6/http';
import { sleep, check } from 'k6';

export let options = {
  vus: 5,
  duration: "30s"
};

export default function() {

  let result = http.get('https://novanetk6api.azurewebsites.net/nocache/99');

  check(result, {
    "Status is 200": (r) => r.status == 200,
    "Duration < 1000ms": (r) => r.timings.duration < 1000
  });

  sleep(1); //sekunder
}
```

# Checks : Resultat

```
✓ Status is 200
X Duration < 1000ms
  ↳ 0% — ✓ 0 / X 70
```

```
checks.....: 50.00% ✓ 70 X 70
data_received.....: 70 kB 2.3 kB/s
data_sent.....: 12 kB 386 B/s
http_req_blocked.....: avg=21.45ms min=0s med=0s max=305.23ms p(90)=0s p(95)=299.22ms
http_req_connecting.....: avg=2.95ms min=0s med=0s max=43ms p(90)=0s p(95)=40.99ms
http_req_duration.....: avg=1.06s min=1.04s med=1.04s max=2.04s p(90)=1.06s p(95)=1.06s
http_req_receiving.....: avg=98.97µs min=0s med=0s max=998.5µs p(90)=777.93µs p(95)=838.47µs
http_req_sending.....: avg=0s min=0s med=0s max=0s p(90)=0s p(95)=0s
http_req_tls_handshaking...: avg=10.19ms min=0s med=0s max=148ms p(90)=0s p(95)=141.09ms
http_req_waiting.....: avg=1.06s min=1.04s med=1.04s max=2.04s p(90)=1.06s p(95)=1.06s
http_reqs.....: 70 2.333333/s
iteration_duration.....: avg=2.08s min=2.04s med=2.04s max=3.34s p(90)=2.06s p(95)=2.36s
iterations.....: 70 2.333333/s
vus.....: 5 min=5 max=5
vus_max.....: 5 min=5 max=5
```



# Thresholds

```
import http from 'k6/http';
import { sleep, check } from 'k6';

export let options = {
  vus: 5,
  duration: "30s",
  thresholds: {
    "errors": ["rate<0.1"], // <10% errors
  }
};

export default function() {

  let result = http.get('https://novanetk6api.azurewebsites.net/nocache/99');

  check(result, {
    "Status is 200": (r) => r.status == 200,
    "Duration < 2000ms": (r) => r.timings.duration < 2000
  });

  sleep(1); //sekunder
}
```

# Tresholds : Resultat

**WARN**[0003] No data generated, because no script iterations finished, consider making the test duration longer

✓ Status is 200  
X Duration < 2000ms  
└ 0% - ✓ 0 / X 4

checks.....	50.00%	✓ 4	X 4						
http_req_blocked.....	avg=310.94ms	min=310.94ms	med=310.94ms	max=310.94ms	p(90)=310.94ms	p(95)=310.94ms			
http_req_connecting.....	avg=39.95ms	min=39.95ms	med=39.95ms	max=39.95ms	p(90)=39.95ms	p(95)=39.95ms			
X http_req_duration.....	avg=1.05s	min=1.05s	med=1.05s	max=1.05s	p(90)=1.05s	p(95)=1.05s			
http_req_receiving.....	avg=0s	min=0s	med=0s	max=0s	p(90)=0s	p(95)=0s			
http_req_sending.....	avg=0s	min=0s	med=0s	max=0s	p(90)=0s	p(95)=0s			
http_req_tls_handshaking...	avg=148.07ms	min=148.07ms	med=148.07ms	max=148.07ms	p(90)=148.07ms	p(95)=148.07ms			
http_req_waiting.....	avg=1.05s	min=1.05s	med=1.05s	max=1.05s	p(90)=1.05s	p(95)=1.05s			
http_reqs.....	4	2.019838/s							
vus.....	5	min=5	max=5						
vus_max.....	5	min=5	max=5						

**ERRO**[0003] some thresholds have failed

# Stages

```
import http from 'k6/http';
import { sleep, check } from 'k6';

export let options = {
  thresholds: {
    "errors": ["rate<0.1"], // <10% errors
  },
  stages: [
    { duration: '10s', target: 10 },
    { duration: '15s', target: 20 },
    { duration: '20s', target: 30 },
    { duration: '30s', target: 100 },
  ],
};

export default function() {
  let result = http.get('https://novanetk6api.azurewebsites.net/nocache/99');

  check(result, {
    "Status is 200": (r) => r.status == 200,
    "Duration < 2000ms": (r) => r.timings.duration < 2000
  });

  sleep(1); //sekunder
}
```

# Sammenligning: Resultat

```
Dictionary Cache duration...: avg=43.3176      min=39.9998      med=42.2671      max=51.2598      p(90)=47.1982      p(95)=49.10078
Lazy Cache duration.....: avg=43.752602     min=40.0301     med=42.9399     max=51.0018     p(90)=48.50005     p(95)=49.69604
Memory Cache duration.....: avg=44.792578     min=40.8522     med=43.0006     max=71.3846     p(90)=49.0873     p(95)=50.86346
No Cache duration.....: avg=1082.556407   min=1042.4406   med=1054.5934   max=1997.289   p(90)=1060.61254   p(95)=1065.64368
Redis Cache duration.....: avg=91.621389     min=76.1802     med=80.01685     max=698.2008     p(90)=85.53435     p(95)=87.9904
checks.....: 100.00% ✓ 574 ✗ 0
data_received.....: 149 kB  5.0 kB/s
data_sent.....: 58 kB  1.9 kB/s
http_req_blocked.....: avg=5.22ms      min=0s      med=0s      max=304.77ms p(90)=0s      p(95)=0s
http_req_connecting.....: avg=747.87µs    min=0s      med=0s      max=46.12ms  p(90)=0s      p(95)=0s
http_req_duration.....: avg=266.76ms    min=39.99ms med=47ms     max=1.99s    p(90)=1.05s   p(95)=1.05s
http_req_receiving.....: avg=134.09µs    min=0s      med=0s      max=7ms      p(90)=931.1µs p(95)=999.67µs
http_req_sending.....: avg=0s          min=0s      med=0s      max=0s      p(90)=0s      p(95)=0s
http_req_tls_handshaking....: avg=2.42ms      min=0s      med=0s      max=145.01ms p(90)=0s      p(95)=0s
http_req_waiting.....: avg=266.63ms    min=39.93ms med=46.99ms  max=1.99s    p(90)=1.05s   p(95)=1.05s
http_reqs.....: 287      9.566361/s
iteration_duration.....: avg=2.58s      min=2.5s      med=2.52s     max=3.47s     p(90)=2.81s     p(95)=2.83s
iterations.....: 54      1.799943/s
vus.....: 5      min=5 max=5
vus_max.....: 5      min=5 max=5
```

# Azure Devops Pipeline

Novanet.K6.Api / azure-pipelines-load-tests.yml

```
# trigger:
# - master

pool:
  vmImage: 'ubuntu-latest'
  **

steps:
- script: |
  sudo apt-key adv --keyserver hkp://keyserver.ubuntu.com:80 --recv-keys 379CE192D401AB61
  echo "deb https://dl.bintray.com/loadimpact/deb-stable-main" | sudo tee -a /etc/apt/sources.list
  sudo apt-get update
  sudo apt-get install k6
  displayName: Install k6 tool

- script: |
  k6 run tests/scripts_99.js
  displayName: Run k6 load test within Azure Pipelines
```



## Jobs in run #20200304.1

Novanet.K6.Api Load Tests (on Build Agent)

✓ Job	1m 14s
✓ Initialize job	1s
✓ Checkout Novanet.K6.Api@maste...	4s
✓ Install k6 tool	37s
✓ Run k6 load test within Azure Pi...	31s
✓ Post-job: Checkout Novanet.K6....	<1s
✓ Finalize Job	<1s
✓ Report build status	<1s

## ✓ Run k6 load test within Azure Pipelines

```
61 time="2020-03-04T11:29:20Z" level=info msg=Running i=48 t=27.983091273s
62 time="2020-03-04T11:29:21Z" level=info msg=Running i=49 t=28.983108919s
63 time="2020-03-04T11:29:22Z" level=info msg=Running i=53 t=29.98312657s
64 time="2020-03-04T11:29:22Z" level=info msg=Test finished" i=53 t=30.000075415s
65
66 ✓ no cache status is 200
67 ✓ memory status is 200
68 ✓ lazy cache status is 200
69 ✓ no cache duration < 2000ms
70 ✓ memory cache duration < 1000m
71 ✓ dictionary status is 200
72 ✓ dictionary cache duration < 1000m
73 ✓ redis status is 200
74 ✓ redis cache duration < 1000m
75 ✓ lazy cache duration < 1000m
76
77 Dictionary Cache duration.... avg=82.291145 min=67.997635 med=68.913962 max=760.660316 p(90)=72.435394 p(95)=73.850587
78 Lazy Cache duration..... avg=85.384457 min=68.104233 med=68.874488 max=763.510992 p(90)=73.551169 p(95)=76.189131
79 Memory Cache duration..... avg=69.891357 min=67.825635 med=68.716752 max=93.252091 p(90)=72.555478 p(95)=73.217008
80 No Cache duration..... avg=1117.574571 min=1069.822308 med=1077.662299 max=1706.743881 p(90)=1085.488822 p(95)=1502.796855
81 Redis Cache duration..... avg=116.438366 min=103.64883 med=104.908577 max=693.939227 p(90)=107.861458 p(95)=108.608465
82 checks..... 100.00% ✓ 536 X 0
83 data_received..... 146 kB 4.9 kB/s
84 data_sent..... 57 kB 1.9 kB/s
85 http_req_blocked..... avg=4.8ms min=1.9µs med=3.9µs max=259.94ms p(90)=7.1µs p(95)=8.1µs
86 http_req_connecting..... avg=1.33ms min=0s med=0s max=72.08ms p(90)=0s p(95)=0s
87 http_req_duration..... avg=295.75ms min=67.82ms med=70.8ms max=1.7s p(90)=1.07s p(95)=1.08s
88 http_req_receiving..... avg=68.8µs min=21.3µs med=54.8µs max=939.42µs p(90)=97.24µs p(95)=109.5µs
89 http_req_sending..... avg=21.61µs min=6.2µs med=19.25µs max=72µs p(90)=37.23µs p(95)=42.85µs
90 http_req_tls_handshaking.... avg=2.58ms min=0s med=0s max=140.93ms p(90)=0s p(95)=0s
91 http_req_waiting..... avg=295.66ms min=67.72ms med=70.71ms max=1.7s p(90)=1.07s p(95)=1.08s
92 http_reqs..... 268 8.933311/s
93 iteration_duration..... avg=2.75s min=2.63s med=2.64s max=4.22s p(90)=2.98s p(95)=3.21s
94 iterations..... 53 1.766662/s
95 vus..... 5 min=5 max=5
96 vus_max..... 5 min=5 max=5
97
98
99
100 Finishing: Run k6 load test within Azure Pipelines
```

# Visualisering

- K6 Cloud (i skyen)
  - SaaS, betalt løsning
  - Testene kjøres i K6 sin sky
  - Får automatisk visualisering og historikk
- InfluxDB og Grafana (on prem eller sky)
  - InfluxDB is the open source time series database
  - Grafana is the open source analytics & monitoring solution for every database.

```
k6 login cloud  
k6 cloud loadtest.js
```

```
docker pull influxdb  
docker run -p 8086:8086 -v c:\temp\influxdb\var\lib\influxdb  
influxdb  
k6 run --out influxdb:http://localhost:8086/ .\loadtest.js
```

```
docker pull grafana/Grafana  
docker run -d --name=grafana -p 3000:3000 grafana/grafana
```

# Ting å tenke på:

- Når bør ytelsestestene kjøre?
  - Manuelt
  - Ved hver release
  - Nattlig
- Hvordan stoppe/rulle tilbake en release når testene feiler?
  - Alert + rulle tilbake manuelt
  - Deployment slots
- Vær snill, ingen DDoS'ing!